

## 0.1 Reversible Transactions over Ethernet Links

Atomic Ethernet supports deterministic reversibility at the transaction level. This enables operations to be undone or rolled back by design—without ambiguity or loss—by embedding invertible transformations into the transmission semantics. Here, we outline a formal foundation for such reversibility using linear algebraic constructs, applied to unidirectional flows over a point-to-point Ethernet link.

GPT-Compressed adaptation of ./AE-Specifications-ETH/chapters/@Grok-reversibility-original.tex

### 0.1.1 Framing Transactions as Linear Operators

Let each transmitted message  $m \in \mathbb{F}_2^n$  be a vector in a binary vector space. The sender applies a reversible transformation  $T \in GL(n, \mathbb{F}_2)$ —an invertible matrix over the Galois field  $\mathbb{F}_2$ —to produce the transmitted payload:

$$m' = T \cdot m$$

At the receiver, the inverse transform  $T^{-1}$  restores the original message:

$$m = T^{-1} \cdot m'$$

This formulation ensures that the transformation is loseless and decodable, and that every transaction can be reversed precisely, assuming both ends share  $T$  and agree on a consistent ordering.

### 0.1.2 Design Implications

Reversibility has several architectural consequences:

- **State Preservation:** Nodes maintain minimal internal state beyond the invertible transformation, supporting rollback without checkpointing.
- **Deterministic Rollback:** If a fault or cancellation occurs, the receiver may return  $m'$  along with  $T^{-1}$  or an encoded undo message, allowing the sender to revert application state.
- **Causal Provenance:** Every transformation  $T$  acts as a causal marker for the transaction's origin and structure. This ensures full verifiability of message lineage and intent.

### 0.1.3 Failure Recovery and Time Symmetry

In conventional systems, rollback is an afterthought—an external protocol layered above an irreversible transmission substrate. By contrast,

Atomic Ethernet supports *built-in reversibility*, which permits time-symmetric protocols where forward progress and backtracking are mirror images.

A link failure mid-transaction results in an incomplete vector transformation. Since the transformation is linear and invertible, partial data receipt can trigger a retry or reversion without ambiguity. The system may encode a rollback transaction as  $-T \cdot m$  or transmit a companion transaction to cancel the original.

#### 0.1.4 Physical and Logical Layer Interface

Reversible transactions reside above the PHY layer, but the encoding scheme must be amenable to in-line streaming and backpressure. Slices are transmitted in fixed-size, entropy-bounded chunks, each marked with the transformation context. Intermediate nodes (in a multi-hop path) may propagate or transform  $T$  according to routing decisions, but the final receiver must be able to apply a composite inverse.

[MISSING FIGURE “linear-reversible”]

#### 0.1.5 Applications and Extensions

This mechanism enables new classes of semantics-aware networking:

- **Reversible Compute Fabric:** Transactions may represent computation steps. Reversal allows rollback of mispredictions or branch divergence in distributed computation.
- **Lossless Failure Handling:** Rather than assuming losses, the protocol treats all disruptions as reversible events, ensuring that no state is committed without an acknowledgment or its inverse.
- **Formal Auditing:** Because each transmission is encoded via known operators, a complete proof-of-history can be generated for compliance or replay.

#### 0.1.6 Mathematical Construction

Let the application payload  $d \in \mathbb{F}_2^n$  represent a fixed-size bit vector, e.g., 512 bits for a standard Ethernet frame. We model  $d$  as a column vector in a binary vector space.

To encode the transaction, the sender selects an invertible matrix  $T \in GL(n, \mathbb{F}_2)$ , producing the transmitted record:

$$r = T \cdot d$$

where:

- $T$  is an  $n \times n$  matrix, generated such that  $\det(T) = 1$ , i.e.,  $T$  is invertible.
- The matrix  $T$  may be predefined, negotiated, or derived from a seed agreed by both parties.
- The receiver recovers  $d$  by applying the inverse:  $d = T^{-1} \cdot r$

*Example:* Suppose  $d = \begin{bmatrix} 1 \\ 0 \\ 1 \\ 1 \end{bmatrix}$  and the transformation matrix is

$$T = \begin{bmatrix} 1 & 1 & 0 & 0 \\ 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 \\ 1 & 0 & 0 & 1 \end{bmatrix}$$

Then the transmitted vector becomes:

$$r = T \cdot d = \begin{bmatrix} 1 \\ 1 \\ 0 \\ 0 \end{bmatrix}$$

This reversible encoding guarantees that no information is lost in transmission, and rollback is always feasible.

### 0.1.7 Encoding Strategies

The matrix  $T$  can be chosen to reflect transaction metadata:

- **Semantic Operators:** Specific rows or structures in  $T$  may denote specific transaction types (e.g., write, abort, checkpoint).
- **Forward-Only Compression:** Even when reversal is not expected, the matrix framework enables low-entropy encoding and structured inference.
- **Replay Protection:** Nonces or keys can be embedded into  $T$  to thwart reapplication of stale transactions.

### 0.1.8 Operational Summary

*Encoding* Sender maps data  $d$  to transmitted record  $r = T \cdot d$

*Transmission* Frame  $r$  is sent with metadata identifying or referencing  $T$

*Reception* Receiver decodes via  $T^{-1}$  to recover  $d$

*Rollback* If reversal is needed, transmit inverse operation using  $T^{-1} \cdot r$  or explicitly send a rollback instruction encoded via a known  $T_{undo}$

This mechanism ensures that every transaction in the link pipeline has a symmetric undo path, forming the basis for reversible computing and auditable networking semantics.

## 0.2 From Clos to Graph: A Shift in Systems Thinking

Conventional datacenter networks built on Clos topologies and Kubernetes orchestration suffer from layered human dependencies: from manual switch configuration to YAML sprawl, from static policy declarations to failure-prone operational recovery. These fragilities compound superlinearly at scale.

Æthernet reimagines this by introducing an alternative: a Direct-Connected 8-Valency IPU mesh, running a decentralized Graph Virtual Machine (GVM). This system inverts operational complexity—delegating routing, scheduling, and security to topology-aware algorithms executed on a uniform graph structure. The goal is zero-trust-by-default, self-partitioning, and human-optional orchestration.

## 0.3 The Graph Virtual Machine

At its core, the GVM exposes programmable primitives for graph manipulation:

- `traverse(type, source, target)` – initiate message routing via BFS, DFS
- `partition(method, subgraphs)` – divide the network using k-means, MST, or spectral cuts
- `optimize(metric, scope)` – apply shortest path, max flow, or latency tuning globally or locally
- `deploy(tenant, subgraph, policy)` – install workloads subject to affinity, proximity, or isolation constraints

Every instruction is compiled into local actions at each node. No global controller is required.

## 0.4 Autonomy at 10,000 Nodes

Unlike Kubernetes, which becomes brittle beyond 1,000 nodes due to centralized control planes, the GVM scales linearly. Each IPU communicates only with 8 neighbors and executes a local GVM. Graph operations (e.g., BFS) run in  $\mathcal{O}(\log n)$  time.

Maintenance, routing, and failure handling are executed as streaming graph updates. No YAML, no kubect!, no overlays.

## 0.5 Security via Confinement and Covers

GVM supports autonomous security through:

- **Graph Covers:** Define disjoint regions of the graph to isolate tenants.
- **Stacked Trees:** Construct independent spanning trees for redundant broadcast or failover paths.
- **Dynamic Partitions:** Tenants exhibiting anomalous behavior are moved to smaller or isolated subgraphs.

Attack surfaces are localized to 32 nodes at most. No blast radius extends across racks.

## 0.6 AI-Augmented Programming

AI assistants integrated with GVM support:

*Policy Translation* Turn human intent into graph ops (e.g., “keep tenant A near B but far from C”).

*Optimization Hints* Suggest `optimize(flow)` or `partition(k-means)` based on real-time metrics.

*Security Monitoring* Isolate compromised nodes via AI-derived instruction chains.

*Graph Debugging* Visualize topologies and bottlenecks, offering suggestions for rerouting or partitioning.

The AI becomes the “network engineer,” shifting the cognitive load away from humans.

## 0.7 Resilience via Topology

The 8-valent IPU mesh has no switches—only direct node-to-node links. This architecture achieves:

- **High Dynamic Laplacian Resilience (DLR):** With 8 neighbors, nodes tolerate up to 7 link failures before isolation.
- **No Leaf Bottlenecks:** Unlike Clos networks, where a failed top-of-rack switch can isolate 32 nodes, this mesh has no such point of failure.
- **Self-Healing:** Graph operations reconfigure routing paths in microseconds.

## 0.8 Mars-Scale Simplicity

In constrained environments—e.g., Martian colonies—human labor is scarce and risk is intolerable. GVM’s autonomy makes it ideal:

- **Resilient:** Survives link loss, power failures, or electromagnetic damage
- **Minimal Ops:** Colonists specify goals; GVM and AI enact them
- **Self-Scaling:** From 100 to 10,000 nodes with no added complexity  
[MISSING FIGURE mars-network.pdf]

## 0.9 Why This Replaces Kubernetes

Metric	Kubernetes/Clos	GVM/IPU
Configuration	Manual YAML, CNI plugins	Self-partitioning graph ops
Operation	Human-tuned schedulers	Fully autonomous
Security	Declarative + fragile	Dynamic + confined
Latency	10–30 $\mu$ s typical	$\sim$ 50–100ns
Failure Domain	Rack-wide (32+ nodes)	1–8 nodes max
Scaling	$n \sim 1,000$ ceiling	$n \geq 10,000$ trivial