

0.1 Mathematical Foundations

In low-latency, high-throughput Layer 2 environments (e.g., Ethernet links), it's useful to model transactions as mathematical operations that can be precisely undone. This enables rollback, audit, and error recovery without heavyweight protocols.

1. Model data as vectors.

Each Ethernet frame is viewed as a vector in $\mathbf{GF}(2)^n$, treating bits not as opaque payload but as elements in a vector space over a finite field.

2. Transactions as invertible operations.

The sender and receiver maintain a shared state $S \in \mathbf{GF}(2)^n$. A transaction is an invertible linear transformation T applied to that state: $S' = T(S)$. Because T is invertible, the original state can always be recovered via T^{-1} .

3. Reversibility via state updates.

To reverse a transaction, one sends a message (or derivable signal) allowing the application of T^{-1} . This guarantees deterministic rollback.

We consider a chain of $N + 1$ nodes labeled $A_0 \rightarrow A_1 \rightarrow \dots \rightarrow A_N$, where each node A_i maintains a local state vector $S_i \in \mathbf{GF}(2)^n$, typically initialized to the all-zero vector 0^n or some other agreed-upon state. Each link ($A_i \rightarrow A_{i+1}$) between adjacent nodes is associated with an invertible linear transformation $T_{i,i+1}$, which governs how state updates propagate along the chain.

0.1.1 Forward Execution

To execute a transaction spanning all links:

1. At each hop i , node A_i applies $T_{i,i+1}$ to its state S_i and transmits the transformation to A_{i+1} .
2. Node A_{i+1} applies the same $T_{i,i+1}$ to its own state S_{i+1} , maintaining link-local consistency.

The result is a chained sequence of transformations:

$$S'_i = T_{i,i+1} \cdot S_i \quad \text{for } i = 0, 1, \dots, N-1, \quad S'_N = T_{N-1,N} \cdot S_N.$$

0.1.2 Rollback (Reverse Direction)

Reversibility is achieved by applying the inverse transformations in reverse order:

1. Node A_N applies $T_{N-1,N}^{-1}$ to revert S'_N to S_N .
2. It signals node A_{N-1} , which applies $T_{N-1,N}^{-1}$ and then $T_{N-2,N-1}^{-1}$, and so on.
3. This continues up the chain until A_0 applies $T_{0,1}^{-1}$, restoring the original S_0 .

Adapted from sections of ./AE-Specifications/chapters/@GPT-reversibility.tex

0.1.3 Example: XOR-Based Masks

If each $T_{i,i+1}$ is a simple XOR with mask $\Delta_{i,i+1}$, then:

$$S_i \mapsto S_i \oplus \Delta_{i,i+1}, \quad S_{i+1} \mapsto S_{i+1} \oplus \Delta_{i,i+1}.$$

Reversing just involves reapplying the same mask due to $\Delta \oplus \Delta = 0$.

0.1.4 Notes on Synchronization

- **Acknowledgments:** Each node should confirm that the next node has applied its transformation before committing its own.
- **Composite View:** The full transaction across N links is a composition:

$$T_{\text{total}} = T_{N-1,N} \circ T_{N-2,N-1} \circ \cdots \circ T_{0,1}.$$

- **Error Handling:** Any failure in transmission or transformation must be detected early, as desynchronization across nodes can compound. Redundant encodings, checksums, or commit/abort protocols may be used.

0.2 Atomic Transactions on \mathcal{AE} -Link

0.2.1 One-Phase Commit

0.2.2 Two-Phase Commit

0.2.3 Four-Phase Commit

0.3 Flow Control and Backpressure

0.4 Transactions on Trees